

```
from google.colab import drive
drive.mount('/content/drive').
```

Run this cell to mount your Google Drive.

DISMISS

browser: https://accounts.google.com/o/oauth2/auth?client_

ation code:

.....

Mounted at /content/drive

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing
import xgboost as xgb
color = sns.color_palette()
```

```
%matplotlib inline
```

```
pd.options.mode.chained_assignment = None # default='warn'
pd.options.display.max_columns = 999
```

```
from subprocess import check_output
#print(check_output(["ls", "../input"]).decode("utf8"))
```

```
train_df = pd.read_csv("/content/drive/My Drive/Projects for Submission/Project 1 - Mercedes-Benz/Train.csv")
test_df = pd.read_csv("/content/drive/My Drive/Projects for Submission/Project 1 - Mercedes-Benz/Test.csv")
print("Train shape : ", train_df.shape)
print("Test shape : ", test_df.shape)
```

```
↳ Train shape : (4209, 378)
   Test shape : (4209, 377)
```

```
train_df.head(,)
```

```
↳
```

	ID	y	x0	x1	x2	x3	x4	x5	x6	x8	x10	x11	x12	x13	x14	x15	x16	x
0	0	130.81	k	v	at	a	d	u	j	o	0	0	0	1	0	0	0	
1	6	88.53	k	t	av	e	d	y	l	o	0	0	0	0	0	0	0	
2	7	76.26	az	w	n	c	d	x	j	x	0	0	0	0	0	0	0	
3	9	80.62	az	t	n	f	d	x	l	e	0	0	0	0	0	0	0	
4	13	78.02	az	v	n	f	d	h	d	n	0	0	0	0	0	0	0	

```
test_df.head(,)
```

```
↳
```

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	X12	X13	X14	X15	X16	X17	X18
0	1	az	v	n	f	d	t	a	w	0	0	0	0	0	0	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	0	0	0	0	0	0	0	0
					f	d	a	j	j	0	0	0	0	1	0	0	0	0
					f	d	z	l	n	0	0	0	0	0	0	0	0	0
4	5	w	s	as	c	d	y	i	m	0	0	0	0	1	0	0	0	0

Run this cell to mount your Google Drive.


DISMISS

```
y = train_df
```

```
train_df.columns
```

```
[>] Index(['Id', 'v2a1', 'hacdor', 'rooms', 'hacapo', 'v14a', 'refrig', 'v18q',
          'v18q1', 'r4h1',
          ...,
          'SQBescolari', 'SQBage', 'SQBhogar_total', 'SQBedjefe', 'SQBhogar_nin',
          'SQBovercrowding', 'SQBdependency', 'SQBmeaned', 'agesq', 'Target'],
          dtype='object', length=143)
```

```
plt.figure(figsize=(8,6))
plt.scatter(range(train_df.shape[0]), np.sort(train_df.y.values))
plt.xlabel('index', fontsize=12)
plt.ylabel('y', fontsize=12)
plt.show()
```



```
ulimit = 180
train_df['y'].ix[train_df['y']>ulimit] = ulimit

plt.figure(figsize=(12,8))
sns.distplot(train_df.y.values, bins=50, kde=False)
plt.xlabel('y value', fontsize=12)
```


```
plt.show().
```

➞ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: .ix is deprecated. Please use .loc for label based indexing or .iloc for positional indexing

Run this cell to mount your Google Drive. [Click here:](#)

DISMISS

<https://pandas-docs/stable/indexing.html#ix-indexer-is-deprec>



```
dtype_df = train_df.dtypes.reset_index()
dtype_df.columns = ["Count", "Column Type"]
dtype_df.groupby("Column Type").aggregate('count').reset_index()
```

➞

	Column Type	Count
0	int64	369
1	float64	1
2	object	8

```
dtype_df.ix[:10,:]
```

➞

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: DeprecationWar
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

Run this cell to mount your Google Drive.

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprec>

launching an IPython kernel.

DISMISS

		type
0	ID	int64
1	y	float64
2	X0	object
3	X1	object
4	X2	object
5	X3	object
6	X4	object
7	X5	object
8	X6	object
9	X8	object
10	X10	int64

```
missing_df = train_df.isnull().sum(axis=0).reset_index()
missing_df.columns = ['column_name', 'missing_count']
missing_df = missing_df.ix[missing_df['missing_count']>0]
missing_df = missing_df.sort_values(by='missing_count')
missing_df
```

```
↳ /usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:3: DeprecationWar
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:

<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprec>

This is separate from the ipykernel package so we can avoid doing imports un

column_name missing_count

```
unique_values_dict = {}
for col in train_df.columns:
    if col not in ["ID", "y", "X0", "X1", "X2", "X3", "X4", "X5", "X6", "X8"]:
        unique_value = str(np.sort(train_df[col].unique()).tolist())
        tlist = unique_values_dict.get(unique_value, [])
        tlist.append(col)
        unique_values_dict[unique_value] = tlist[:]
for unique_val, columns in unique_values_dict.items():
    print("Columns containing the unique values : ", unique_val)
    print(columns)
```

```
print(, "-----").
```

```

Columns containing the unique values :  [0, 1]
['x10', 'x12', 'x13', 'x14', 'x15', 'x16', 'x17', 'x18', 'x19', 'x20', 'x21',
-----
Columns containing the unique values :  [0]
', 'x233', 'x235', 'x268', 'x289', 'x290', 'x293', 'x297',
-----

```


Run this cell to mount your Google Drive.

DISMISS

```

var_name = "X0"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.stripplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show()

```




```

var_name = "X1"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.stripplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show()

```

↳

Distribution of y variable with X1




Run this cell to mount your Google Drive.

DISMISS

```
var_name = "X2"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.boxplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show()
```



Distribution of y variable with X2



```
var_name = "X3"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.violinplot(x=var_name, y='y', data=train_df, order=col_order)
```


```
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show(.,)
```



Distribution of y variable with X3

Run this cell to mount your Google Drive.


DISMISS



```
var_name = "X4"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.violinplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show(.,)
```



Distribution of y variable with X4




```

var_name = "X5"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.boxplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)

```

Run this cell to mount your Google Drive.

DISMISS




```

var_name = "X6"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.boxplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show()

```



Distribution of y variable with X6




Run this cell to mount your Google Drive.

DISMISS

```
var_name = "X8"
col_order = np.sort(train_df[var_name].unique()).tolist()
plt.figure(figsize=(12,6))
sns.boxplot(x=var_name, y='y', data=train_df, order=col_order)
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show()
```



Distribution of y variable with X8



```
zero_count_list = []
one_count_list = []
cols_list = unique_values_dict['[0, 1]']
```

```
for col in cols_list:
    zero_count_list.append((train_df[col]==0).sum())
    one_count_list.append((train_df[col]==1).sum())

N = len(cols_list)
ind = np.arange(N)
width = 0.35
```

Run this cell to mount your Google Drive.

DISMISS


```
)
    unt_list, width, color='red')
    nt_list, width, left=zero_count_list, color="blue")

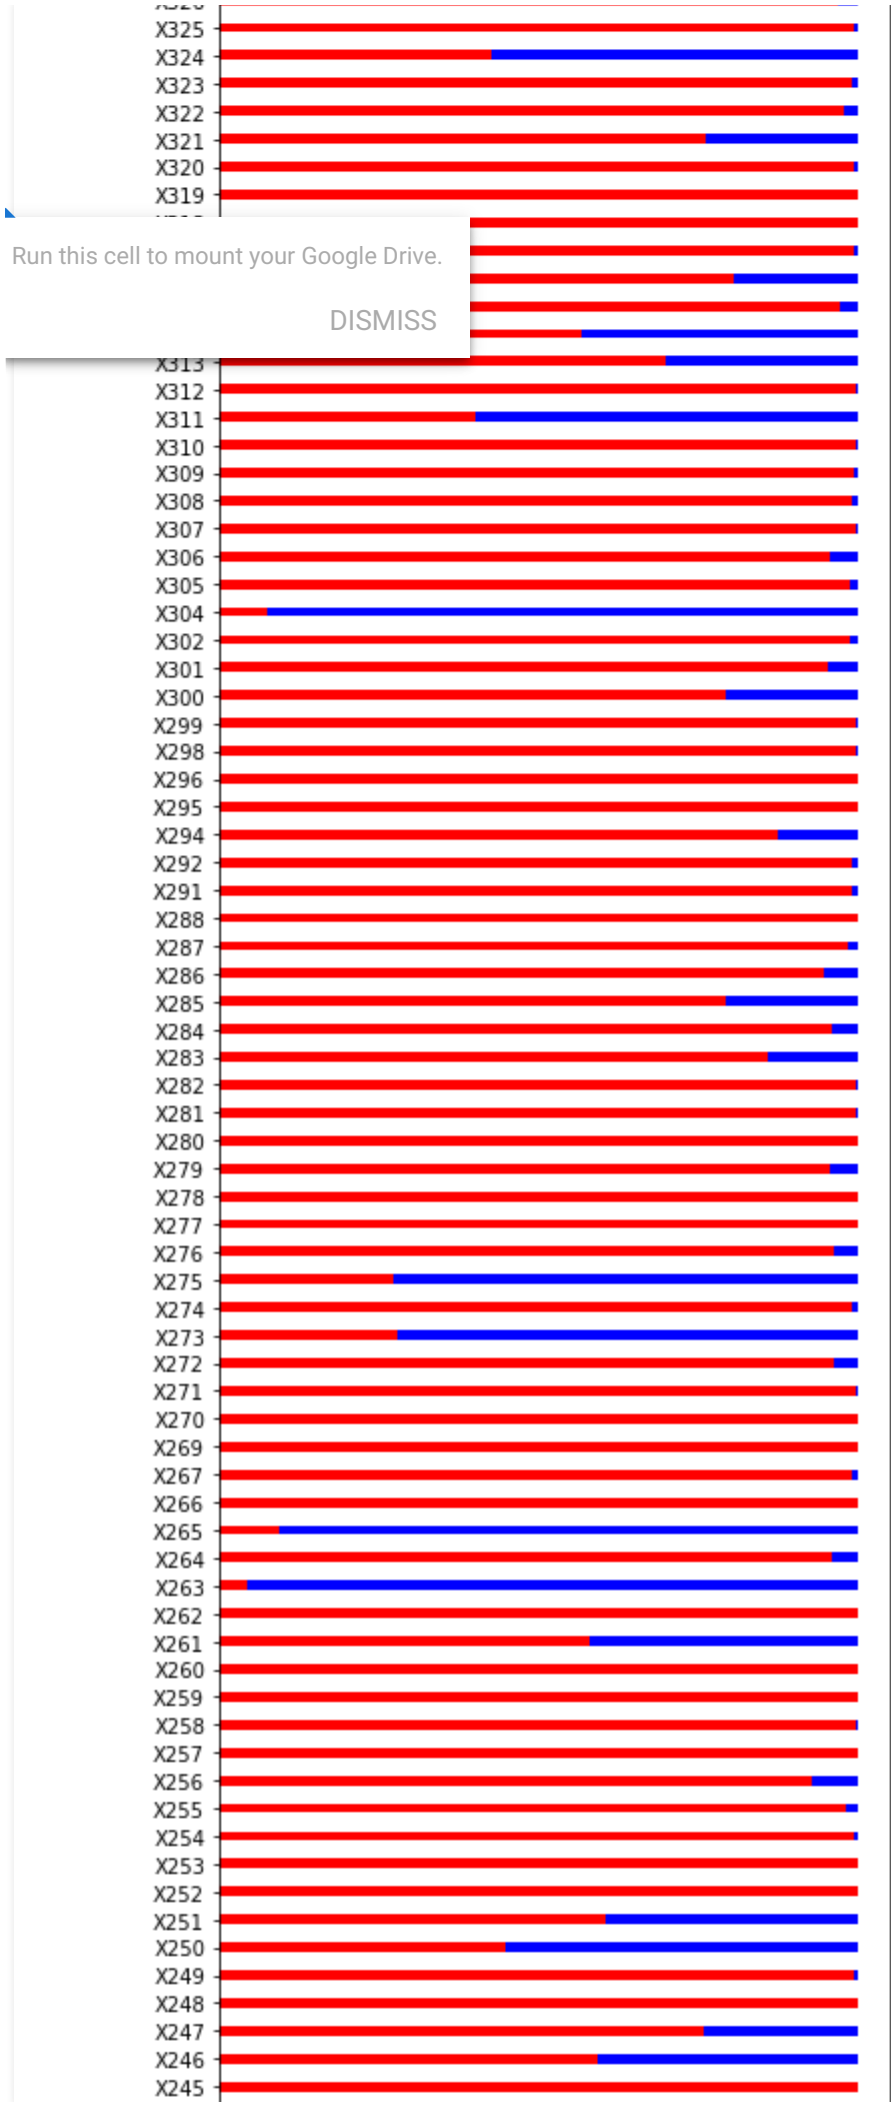
plt.legend((p1[0], p2[0]), ('Zero count', 'One Count'))
plt.show()
```




Run this cell to mount your Google Drive.

DISMISS











Run this cell to mount your Google Drive.


DISMISS





Run this cell to mount your Google Drive.

DISMISS



Run this cell to mount your Google Drive.

DISMISS

0 500 1000 1500 2000 2500 3000 3500 4000

```
zero_mean_list = []
one_mean_list = []
cols_list = unique_values_dict['[0, 1]']
for col in cols_list:
    zero_mean_list.append(train_df.ix[train_df[col]==0].y.mean())
    one_mean_list.append(train_df.ix[train_df[col]==1].y.mean())

new_df = pd.DataFrame({"column_name":cols_list+cols_list, "value":[0]*len(cols_list) + [1]*len(cols_list)})
new_df = new_df.pivot('column_name', 'value', 'y_mean')

plt.figure(figsize=(8,80))
sns.heatmap(new_df)
plt.title("Mean of y value across binary variables", fontsize=15)
plt.show()
```




```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:5: DeprecationWar
.ix is deprecated. Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:


Run this cell to mount your Google Drive. <http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprec>


DISMISS

```
on3.6/dist-packages/ipykernel_launcher.py:6: DeprecationWar
Please use
.loc for label based indexing or
.iloc for positional indexing
```

See the documentation here:


<http://pandas.pydata.org/pandas-docs/stable/indexing.html#ix-indexer-is-deprec>





Run this cell to mount your Google Drive.


DISMISS






Run this cell to mount your Google Drive.

DISMISS






Run this cell to mount your Google Drive.

DISMISS

```
var_name = "ID"
plt.figure(figsize=(12,6))
sns.regplot(x=var_name, y='y', data=train_df, scatter_kws={'alpha':0.5, 's':30})
plt.xlabel(var_name, fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of y variable with "+var_name, fontsize=15)
plt.show(.)
```



Distribution of y variable with ID




```
plt.figure(figsize=(6,10))
train_df['eval_set'] = "train"
test_df['eval_set'] = "test"
full_df = pd.concat([train_df[["ID", "eval_set"]], test_df[["ID", "eval_set"]]], axis=0)

plt.figure(figsize=(12,6))
sns.violinplot(x="eval_set", y='ID', data=full_df)
plt.xlabel("eval_set", fontsize=12)
plt.ylabel('y', fontsize=12)
plt.title("Distribution of ID variable with evaluation set", fontsize=15)
plt.show()
```

☞ <Figure size 432x720 with 0 Axes>

Distribution of ID variable with evaluation set



```

for f in ["X0", "X1", "X2", "X3", "X4", "X5", "X6", "X8"]:
    lbl = preprocessing.LabelEncoder()
    lbl.fit(list(train_df[f].values))
    train_df[f] = lbl.transform(list(train_df[f].values))

train_y = train_df['y'].values
train_X = train_df.drop(["ID", "y", "eval_set"], axis=1)

```

Run this cell to mount your Google Drive.

DISMISS

```

s #
rain):
bel()
labels, preds)

```

```


xgb_params = {
    'eta': 0.05,
    'max_depth': 6,
    'subsample': 0.7,
    'colsample_bytree': 0.7,
    'objective': 'reg:linear',
    'silent': 1
}
dtrain = xgb.DMatrix(train_X, train_y, feature_names=train_X.columns.values)
model = xgb.train(dict(xgb_params, silent=0), dtrain, num_boost_round=100, feval=xgb_r2_s

# plot the important features #
fig, ax = plt.subplots(figsize=(12,18))
xgb.plot_importance(model, max_num_features=50, height=0.8, ax=ax)
plt.show()

```



[06:18:03] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear




```
from sklearn import ensemble
model = ensemble.RandomForestRegressor(n_estimators=200, max_depth=10, min_samples_leaf=4)
model.fit(train X, train y)
importances = model.feature_importances_.values
```

Run this cell to mount your Google Drive.


DISMISS

```
indices = np.argsort(importances)[::-1][:20]

plt.figure(figsize=(12,12))
plt.title("Feature importances")
plt.bar(range(len(indices)), importances[indices], color="r", align="center")
plt.xticks(range(len(indices)), feat_names[indices], rotation='vertical')
plt.xlim([-1, len(indices)])
plt.show()
```



Feature importances



Run this cell to mount your Google Drive.

DISMISS